

```
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.callbacks import ReduceLROnPlateau
import matplotlib.pyplot as plt
import tensorflow_datasets as tfds
```

```
dataset, info = tfds.load('tf_flowers', split='train', with_info=True)
```

Downloading and preparing dataset 218.21 MiB (download: 218.21 MiB, generated: 221.83 MiB, total: 440.05 MiB) [00:04<00:00, 1.13 file/s]

Dataset tf_flowers downloaded and prepared to /root/tensorflow_datasets/tf_flowers/3.0.1. Subsequent calls will reuse this data.

```
def preprocess(example):
    image = example['image']
    label = example['label']
    image = tf.image.resize(image, (150, 150))
    image = tf.cast(image, tf.float32) / 255.0
    return image, label
```

```
dataset = dataset.map(preprocess)
```

```
train_size = info.splits['train'].num_examples
train_dataset = dataset.take(train_size * 80 // 100)
test_dataset = dataset.skip(train_size * 80 // 100)
```

```
train_dataset = train_dataset.shuffle(1000).batch(32)
test_dataset = test_dataset.batch(32)
```

```
for image, label in train_dataset.take(1):
    plt.figure()
    plt.imshow(image[0])
    plt.title('Sample Image')
    plt.axis('off')
    plt.show()
    break
```



Sample Image



```
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
```

```

layers.Conv2D(128, (3, 3), activation='relu'),
layers.MaxPooling2D((2, 2)),
layers.Flatten(),
layers.Dense(512, activation='relu'),
layers.Dense(5, activation='softmax')
])

```

⚡ /usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

```

```
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=5, min_lr=0.001)
```

```

history = model.fit(train_dataset, epochs=5, validation_data=test_dataset,
                    callbacks=[reduce_lr])

```

⚡ Epoch 1/5
92/92 ————— **185s** 2s/step - accuracy: 0.3728 - loss: 1.4281 - val_accuracy: 0.4796 - val_lo
Epoch 2/5
92/92 ————— **185s** 2s/step - accuracy: 0.5572 - loss: 1.0681 - val_accuracy: 0.5763 - val_lo
Epoch 3/5
92/92 ————— **203s** 2s/step - accuracy: 0.6394 - loss: 0.8924 - val_accuracy: 0.6485 - val_lo
Epoch 4/5
92/92 ————— **193s** 2s/step - accuracy: 0.6783 - loss: 0.8299 - val_accuracy: 0.6866 - val_lo
Epoch 5/5
92/92 ————— **207s** 2s/step - accuracy: 0.7144 - loss: 0.6984 - val_accuracy: 0.6962 - val_lo

```
plt.figure(figsize=(12, 5))
```

```

# Accuracy
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

```

```

# Loss
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

```

```

plt.tight_layout()
plt.show()

```

